

Rust Memory Management

Kenjiro Taura

2025/07/07

Contents

Introduction 2

Rust Basics 6

Owning Pointers 10


Box<T> type 15

Borrowing pointers (&T) 24

Borrow checking details 28

Summary 56


Introduction



Rust's basic idea to memory management

- Rust maintains that, for any live object,
 1. there is *one and only one* pointer that “owns” it (*the owning pointer*)
 2. there are any number of non-owning pointers to it (*borrowing pointers*)
 3. *borrowing pointers cannot be dereferenced after the owning pointer goes away*
- *⇒ it can safely reclaim the data when the owning pointer goes away*

“*single-ownership* rule”



```
{  
  let a = S{x: ..., y: ...};  
  ...  
} // what a points to will be gone here
```

The rules are enforced statically

- Rust enforces the rules (or, detect violations thereof)
 - *statically*, not *dynamically*
 - *compile-time*, not at *runtime*
 - *before* execution, not *during* execution

“borrow checker”

Escaping from the single ownership model

- there are actually some ways to get around the rules
 1. **reference counting pointers** (\approx multiple owning pointers)
 - counts the number of owners *at runtime*, and reclaim the data when all these pointers are gone
 2. **unsafe/raw pointers** (\approx totally up to you)

they are not specific to Rust, and we'll not cover them below

Rust Basics




Pointer-like data types in Rust

given a type T (i32, struct, enum, ...), below are types representing “references (pointers) to T ”

1. T : **owning** pointer to T
2. $\text{Box}<T>$ (pronounced “box T ”) : **owning** pointer to T
3. $\&T$ (pronounced “ref T ”) : **borrowing pointer** to T
4. $\text{Rc}<T>$ and $\text{Arc}<T>$: shared (reference-counting) owning pointer to T
5. $*T$: unsafe pointer to T

following discussions are focused on
 T , $\text{Box}<T>$ and $\&T$



Pointer-making expressions

given an expression e of type T , below are expressions that make pointers to the value of e (besides e itself)

- `Box::new(e)` (of type `Box< T >`) : an owning pointer
- `& e` (of type `& T`) : a borrowing pointer

An example

```
{  
  let a: S = S{x: ...};           // allocate memory for S  
                                   // and make an owning pointer to it  
  let b: S = a                   // an owning pointer  
  let c: Box<S> = Box::<S>::new(a) // an owning pointer  
  let d: &S = &a                 // a borrowing pointer  
}
```

- note: type of variables can be omitted (spelled out for clarity)
- note: the above program violates several rules so it does not compile

Owning Pointers


Assignments of owning pointers

- to maintain the “single-owner” rule, an assignment of owning pointers in Rust *does not copy, but moves it* out of the righthand side, disallowing further use of it

```
b = a; // a cannot be used below
```

```
fn foo() {  
    let a = S{x: ..., y: ...};  
    ... a.x ...; // OK, as expected  
    ... a.y ...; // OK, as expected
```

```
}
```




Assignments of owning pointers

- to maintain the “single-owner” rule, an assignment of owning pointers in Rust *does not copy, but moves it* out of the righthand side, disallowing further use of it

```
b = a; // a cannot be used below
```


```
fn foo() {  
    let a = S{x: ..., y: ...};  
    ... a.x ...; // OK, as expected  
    ... a.y ...; // OK, as expected  
    // the reference moves out from a  
    let b = a;  
    a.x; // NG, the value has moved out  
    b.x; // OK  
}
```



Argument-passing also moves the reference

- passing a value to a function also moves the reference out of the source


```
fn foo() {  
  let a = S{x: ..., y: ...};  
  ... a.x ...; // OK, as expected  
  ... a.y ...; // OK, as expected  
  
}
```



Argument-passing also moves the reference

- passing a value to a function also moves the reference out of the source

```
fn foo() {  
  let a = S{x: ..., y: ...};  
  ... a.x ...; // OK, as expected  
  ... a.y ...; // OK, as expected  
  // moves the reference out of a  
  f(a);  
  a.x; // NG, the reference has moved  
}
```



Exceptions to “assignment moves the reference”

- you may notice the moving assignment contradicts what you have seen

```
b = a; // a cannot be used after this
```

- if it applies everywhere, does the following program violate it?

```
fn foo() -> f64 {  
    let a = 123.456;  
    let b = a; // does the reference to 123.456 move out from a!?  
    a + 0.789 // if so, is this invalid!?  
}
```

- answer: no, it does *not* apply to primitive types like i32, f64, etc.
- more generally, it does not apply to data types that implement *Copy* trait


Copy trait

- define your struct with `#[derive(Copy, Clone)]` like

```
#[derive(Copy, Clone)]  
struct S { ... }
```

- \Rightarrow assignment or argument-passing of *S* copies the righthand side

```
fn foo() {  
    let a = S{x: ..., y: ...};  
    a.x; // OK, as expected  
    a.y; // OK, as expected
```



```
}
```

- note: copy types trivially maintain the single-owner rule


Copy trait

- define your struct with `#[derive(Copy, Clone)]` like

```
#[derive(Copy, Clone)]  
struct S { ... }
```

- \Rightarrow assignment or argument-passing of *S* copies the righthand side

```
fn foo() {  
    let a = S{x: ..., y: ...};  
    a.x; // OK, as expected  
    a.y; // OK, as expected  
    // the value is copied  
    let b = a;  
    a.x; // OK  
    b.x; // OK, too  
}
```




- note: copy types trivially maintain the single-owner rule

Box<*T*> type

Box<T> makes an owning pointer

- making a pointer by `Box::new(v)` moves the reference out of v , too, and `Box::new(v)` becomes the owning pointer


```
fn foo() {  
    let a = S{x: ..., y: ...};  
    a.x; // OK, as expected  
    a.y; // OK, as expected  
}
```



Box<T> makes an owning pointer

- making a pointer by `Box::new(v)` moves the reference out of `v`, too, and `Box::new(v)` becomes the owning pointer

```
fn foo() {  
  let a = S{x: ..., y: ...};  
  a.x; // OK, as expected  
  a.y; // OK, as expected  
  // OK, now b is the owning pointer  
  let b = Box::new(a)  
  a.x; // NG, the value has moved out  
  (*b).x; // OK  
  b.x; // OK. abbreviation of (*b).x  
}
```



Difference between T and `Box<T>`?

- as you have seen, the effects of

```
let b = a
```

and



```
let b = Box::new(a)
```

look very similar (identical)

- as far as data lifetime is concerned, it is in fact safe to say they are
- Rust has distinction between them for
 1. specifying data layout
 2. allowing dynamic dispatch only for `Box<T>`
 3. specifying where data are allocated (stack vs. heap)

Data layout differences between T and $\text{Box}\langle T \rangle$

- S and U below have different data layouts
 - struct S { ..., p: T , } “embeds” a T into S
 - struct U { ..., p: $\text{Box}\langle T \rangle$, } has p point to a separately allocated T





Data layout differences between T and $\text{Box}\langle T \rangle$

- in particular, $\text{Box}\langle T \rangle$ is essential to define recursive data structures
 - `struct S { ..., p: S, }` is not allowed, whereas
 - `struct U { ..., p: Box<U>, }` is
- note: U above can never be constructed; a recursive data structure typically looks like
 - `struct U { ..., p: Option<Box<U>>, }`

Data layout differences between T and $\text{Box}<T>$

- the distinction is insignificant when discussing lifetimes



- in both cases, data of T (yellow box) is gone exactly when the enclosing structure is gone
- another difference is that Rust allocates T on stack and move it to heap when $\text{Box}<T>$ is made
- again, it has nothing to do with lifetime (unlike C/C++)



Owning pointers and control flows

- Rust compiler determines, for each variable of owning pointer type (T or $\text{Box}<T>$), at which point the variable is can be used or *valid* (i.e., the value has not been moved out)
- it may be a *conservative* estimate

```
fn foo() {  
    let a = S{x: ..., y: ...};  
    if ... {  
        let b = a;  
    }  
    ... a.x ... // NG  
}  
  
fn foo() {  
    let a = S{x: ..., y: ...};  
    for ... {  
        let b = a; // NG  
    }  
}
```

A (huge) implication of the single-owner rule

- with only owning pointers (T and $\text{Box}\langle T \rangle$),
 - you can make *a tree* of data,
 - but you *cannot make a general graph* with joins or cycles, where *a node may be pointed to by multiple nodes*
- to make a graph whose nodes are T , use either
 - $\&T$ to represent edges, or
 - $\text{Vec}\langle T \rangle$ to represent nodes and $\text{Vec}\langle (\text{i32}, \text{i32}) \rangle$ to represent edges




The (huge) implication to memory management

- with only owning pointers (i.e., no borrowing pointers)
- *whenever an owning pointer is gone* (e.g.,
 - ▶ a variable goes out of scope or
 - ▶ a variable or field is overwritten),

the entire tree rooted from the pointer can be safely reclaimed

- Rust exactly does that, with the additional guarantee that *borrowing pointers are never dereferenced after its owning pointer is gone*



Borrowing pointers ($\&T$)

- you can derive any number of borrowing pointers ($\&T$) from T or $\text{Box}\langle T \rangle$
- the owning pointer remains valid after a borrowing pointer has been made

```
let a = S{x: .., y: ..};  
let b = &a;  
... a.x + b.x ... // OK
```

- the issue is how to prevent a program from *dereferencing borrowing pointers after its owning pointer is gone*

Borrowers rule in action

- a borrowing pointer cannot be dereferenced after its owning pointer is gone

```
fn foo() -> i32 {  
    let c: &S; // a reference to S
```

c : &S

```
}
```

Borrowers rule in action

- a borrowing pointer cannot be dereferenced after its owning pointer is gone

```
fn foo() -> i32 {  
    let c: &S; // a reference to S  
    { // an inner block  
        let b: &S; // another reference  
  
    }  
}
```

c : &S

b : &S

Borrowers rule in action


- a borrowing pointer cannot be dereferenced after its owning pointer is gone

```
fn foo() -> i32 {  
  let c: &S; // a reference to S  
  { // an inner block  
    let b: &S; // another reference  
    let a = S{x: ...}; // allocate S  
  
  }  
}
```

c : &S

b : &S


a : S



Borrowers rule in action

- a borrowing pointer cannot be dereferenced after its owning pointer is gone


```
fn foo() -> i32 {  
  let c: &S; // a reference to S  
  { // an inner block  
    let b: &S; // another reference  
    let a = S{x: ...}; // allocate S  
    // OK (both a and b live only until the end of  
    the inner block)  
    b = &a;  
  }  
}
```



Borrowers rule in action

- a borrowing pointer cannot be dereferenced after its owning pointer is gone


```
fn foo() -> i32 {  
  let c: &S; // a reference to S  
  { // an inner block  
    let b: &S; // another reference  
    let a = S{x: ...}; // allocate S  
    // OK (both a and b live only until the end of  
    // the inner block)  
    b = &a;  
    c = b; // dangerous (c outlives a)  
  }  
}
```



Borrowers rule in action

- a borrowing pointer cannot be dereferenced after its owning pointer is gone

```
fn foo() -> i32 {  
  let c: &S; // a reference to S  
  { // an inner block  
    let b: &S; // another reference  
    let a = S{x: ...}; // allocate S  
    // OK (both a and b live only until the end of  
    // the inner block)  
    b = &a;  
    c = b; // dangerous (c outlives a)  
  } // a dies here, making c a dangling pointer  
  c.x // NG (deref a dangling pointer)  
}
```



A *mutable* borrowing reference (&mut T)

- data cannot be modified through ordinary borrowing references & T

```
let a : S = S{x: 10, y: 20};  
let b : &S = &a;  
b.x = 100; // NG
```

- i.e., & T is the type of *immutable* references
- you can modify data only through *a mutable reference* (&mut T)

```
let mut a : S = S{x: 10, y: 20};  
let b : &mut S = &mut a;  
b.x = 100; // OK
```

- the difference is largely orthogonal to memory management

Borrow checking details

A technical remark about the borrow checking

- it's *not a creation* of a dangling pointer, *per se*, that is not allowed, but *dereferencing* of it
- *a slightly modified code below compiles without an error*, despite that `c` becomes a dangling pointer to `a` (as it is not dereferenced past `a`'s lifetime)

```
fn foo() -> i32 {  
    let c: &S; // a reference to S  
    { // an inner block  
        let b: &S; // another reference  
        let a = S{x: ...}; // allocate S  
        // OK (both a and b live only until the end of the inner block)  
        b = &a;  
        c = b; // dangerous (c outlives a)  
    } // a dies here, making c a dangling pointer  
    // c.x don't deref c }  
}
```

How borrow checking works : lifetime

- *lifetime* of data
 - = program points where the data has not been deallocated
 - = program points where the data's owning pointer is valid
- for each borrowing pointer, Rust compiler determines the *lifetime* of data it points to (*referent lifetime*) as its static type
- upon assignment $p = q$ between borrowing pointers, it demands
referent lifetime of $p \subset$ referent lifetime of q

How borrow checking basically works

```
fn foo() -> i32 {  
  let c: &S;  
  {  
    let b: &S;  
    let a = S{x: ...}; // lives until  $\alpha$   
    b = &a; // b's referent lifetime  $\subset$  a's =  $\alpha$   
    c = b; // c's referent lifetime  $\subset$  b's =  $\alpha$   
  } // a dies here ( $\alpha$ )  
  c.x // NG (deref outside c's referent lifetime =  $\alpha$ )  
}
```

1. the owning pointer a's lifetime is the inner block (call it α)
2. due to the assignments,
 - b's referent lifetime $\subset \alpha$
 - c's referent lifetime $\subset \alpha$
3. \therefore c.x outside the inner block ($\not\subset \alpha$) is invalid

Programming with borrowing references

- in more general cases, programs using borrowing references must help compilers track their referent lifetimes
- this must be done for functions called from unknown places, function calls to unknown functions and data structures
- to this end, the programmer sometimes must annotate *reference types with their referent lifetimes*

References in function parameters

- problem: how to check the validity of:

1. functions taking references, *without knowing all its callers*,

```
fn p_points_q(p: &mut P, q: &Q) {  
    p.x = q; //OK?  
}
```

2. and function calls passing references, *without knowing the definition of f?*

```
let c = ...;  
{  
    let a = Q{...};  
    let b = &a;  
    f(c, b);  
} ... c.x.y ... //OK?
```

References in data structures

- problem: how to check the validity of:


1. dereferencing a pointer obtained from a data structure

```
fn ref_from_struct() {  
    ...  
    let p: &P = a.p;  
    ...  
    p.x // OK?  
}
```

2. what about functions taking data structures containing references and returning another containing references, etc.?

Reference type with a lifetime parameter


- to address this problem, Rust's borrowing reference types ($\&T$ or $\&\text{mut } T$) carry *lifetime parameter representing their referent lifetimes*



lives until 'a


Reference type with a lifetime parameter

- to address this problem, Rust's borrowing reference types ($\&T$ or $\&\text{mut } T$) carry *lifetime parameter representing their referent lifetimes*
- syntax:
 - $\&'a T$: reference to “ T whose lifetime is $'a$ ”
 - $\&'a \text{mut } T$: ditto; except you can modify data through it



Reference type with a lifetime parameter

- to address this problem, Rust's borrowing reference types ($\&T$ or $\&\text{mut } T$) carry *lifetime parameter representing their referent lifetimes*
- syntax:
 - $\&'a T$: reference to “ T whose lifetime is ‘ a ”
 - $\&'a \text{mut } T$: ditto; except you can modify data through it
- *every* reference carries a lifetime parameter, though there are places you can omit them
- roughly, you must write them explicitly in function parameters, return types, and struct/enum fields; and can omit them for local variables



Attaching lifetime parameters to functions

- the following does not compile:

```
fn foo(ra: &i32, rb: &i32, rc: &i32) -> &i32 {  
    ra  
}
```

with errors like:

```
|  
| fn foo(ra: &i32, rb: &i32, rc: &i32) -> &i32 {  
|          ----      ----      ----      ^ expected named lifetime parameter  
|  
= help: this function's return type contains a borrowed value, but the signature does not say  
whether it is borrowed from `ra`, `rb`, or `rc`  
help: consider introducing a named lifetime parameter  
|  
| fn foo<'a>(ra: &'a i32, rb: &'a i32, rc: &'a i32) -> &'a i32 {  
|          ++++      ++      ++      ++      ++  
|
```


Why do we need an annotation, *fundamentally*?

- without any annotation, how to know whether this is safe, *without knowing the definition of foo?*

```
{  
  let r : &i32;  
  let a = 123;  
  {  
    let b = 456;  
    {  
      let c = 789;  
      r = foo(&a, &b, &c);  
    }  
  }  
  *r  
}
```

- essentially, the compiler complains “tell me what kind of lifetime foo(&a, &b, &c) has”

Attaching lifetime parameters to functions

- syntax:

```
fn f<'a, 'b, 'c, ...>(p0 : T0, p1 : T1, ...) -> Tr {  
    ...  
}
```

T_0, T_1, \dots , and T_r may use 'a, 'b, 'c, ... as lifetime parameters (e.g., &'a i32)

- $f<'a, 'b, 'c, \dots>$ is a function that takes parameters of respective lifetimes

One way to attach lifetime parameters to the example

```
fn foo<'a>(ra: &'a i32, rb: &'a i32, rc: &'a i32) -> &'a i32
```

- effect: the return value is assumed to point to the shortest of the three
- why? generally, when Rust compiler finds `foo(x, y, z)`, it tries to determine 'a so that 'a \subset referent lifetimes of x, y, and z
- in this case,

'a \subset (life time of a) \cap (life time of b) \cap (life time of c)

- as a result, our program does not compile, even if `foo(&a, &b, &c)` in fact returns &a

```
{
  let r: &i32;
  let a = 123;
  {
    let b = 456;
    {
      let c = 789;
      r = foo(&a, &b, &c);
      // 'a  $\leftarrow \alpha \cap \beta \cap \gamma = \gamma$ 
      // and r's type becomes  $\&\gamma$  i32
    } // c's lifetime (=  $\gamma$ ) ends here
  } // b's lifetime (=  $\beta$ ) ends here
  *r // NG, as we are outside  $\gamma$ 
} // a's lifetime (=  $\alpha$ ) ends here
```

An annotation that works

```
fn foo<'a, 'b, 'c>(ra: &'a i32, rb: &'b i32, rc: &'c i32) -> &'a i32
```

- signifies that the return value points to data whose lifetime is `ra`'s referent lifetime (and has nothing to do with `rb`'s or `rc`'s)
- for `foo(x, y, z)`, Rust compiler tries to determine `'a` so that `'a \subset referent lifetimes of x`
- as a result, the program we are discussing compiles:

```
{
  let r: &i32;
  let a = 123;
  {
    let b = 456;
    {
      let c = 789;
      r = foo(&a, &b, &c);
      // 'a  $\leftarrow \alpha$ 
      // and r's type becomes & $\alpha$  i32
    } // c's lifetime (=  $\gamma$ ) ends here
  } // b's lifetime (=  $\beta$ ) ends here
  *r // OK, as here is within  $\alpha$ 
} // a's lifetime (=  $\alpha$ ) ends here
```

Types with lifetime parameters capture/constrain the function's behavior

- what if you try to fool the compiler by:

```
fn foo<'a, 'b, 'c>(ra: &'a i32, rb: &'b i32, rc: &'c i32) -> &'a i32 {  
    rb  
}
```

- the compiler rejects returning `rb` (of type `&'b`) when the function's return type is `&'a`
- in general, *the compiler allows assignments only between references having the same lifetime parameter*

Another example (make a reference between inputs)

- what if we rewrite

```
r = foo(&a, &b, &c);
```

into

```
bar(&mut r, &a, &b, &c);
```

with bar something like

```
fn bar(r: &mut &i32, a: &i32, b: &i32, c: &i32) {  
    *r = a;  
}
```

Make a reference between inputs

- how to specify lifetime parameters so that:
 1. `*r = a`; in `bar`'s definition is allowed, and
 2. we can dereference `*r` at the end of the caller?

```
{  
  let a = 123;  
  let mut r = &0;  
  {  
    let b = 456;  
    {  
      let c = 789;  
      bar(&mut r, &a, &b, &c); // r → ???  
    } // c's lifetime (=  $\gamma$ ) ends here  
  } // b's lifetime (=  $\beta$ ) ends here  
  *r // OK???  
} // a's lifetime (=  $\alpha$ ) ends here
```

- again, we need to signify that `r` points to `a` (and not `b` or `c` after `bar(&r, &a, &b, &c)`)
- a working lifetime parameter is the following:

```
fn bar<'a, 'b, 'c>(r: &mut &'a i32,  
                  a: &'a i32, b: &'b i32, c: &'c i32) {  
    *r = a;  
}
```


References in data structures

- problem: how to check the validity of programs using data structure containing a borrowing reference

```
struct R {  
  p: &i32  
  ...  
}
```

and functions returning R:

```
fn ret_r(a: &i32, b: &i32, c: &i32) -> R {  
  R { p: a }  
}
```

or taking R (or reference to it):

References in data structures

```
fn take_r(r: &mut R, a: &i32, b: &i32, c: &i32) {  
    r.p = a;  
}
```

References in data structures

- you cannot simply have a field of type $\&T$ in struct/enum like this:

```
struct R {  
    p: &i32  
    ...  
}
```

- you need to specify the lifetime parameter of p , and let R take the lifetime parameter:

```
struct R<'a> {  
    p: &'a i32  
    ...  
}
```

- $R<'a>$ represents *R whose p field points to an $i32$ whose lifetime is $'a$*

References in data structures

- this way, a structure containing borrowing references exposes their referent lifetimes to its user

Attaching lifetime parameters to data structure

- say we like to have data structures:

```
struct T { x: i32 }  
struct S { p: &T }
```

- and a function:

```
fn make_s(a: &T, b: &T) -> S { S { p: a } }
```

- so that the following compiles:

```
let s;  
let a = T{...};  
{  
  let b = T{...};  
  s = make_s(&a, &b);  
}
```

Attaching lifetime parameters to data structure

```
}  
s.p.x
```

- the compiler needs to verify `s.p` points to `a`, not `b`
- we have to signify that by appropriate lifetime parameters

- define `S<'a>` so its `p`'s referent lifetime is `'a`

```
struct S<'a> { p: &'a T }
```

- define `make_s` so it returns `S<'a>` where `'a` is the referent lifetime of its *first* parameter

```
fn make_s(a: &'a T, b: &'b T) -> S<'a> {  
    S { p: a }  
}
```

A more complex example Rust cannot verify

- say we now have data structures

```
struct T { x: i32 }  
struct S {  
    p: &T,  
    q: &T  
}
```

- and a function

```
fn make_s(a: &T, b: &T) -> S {  
    S { p: a, q: b }  
}
```

- so that the following compiles

A more complex example Rust cannot verify

```
let s;  
let a = T{...};  
{  
    let b = T{...};  
    s = make_s(&a, &b);  
}  
s.p.x
```

- again, the compiler needs to verify `s.p` points to `a`, not `b`

Answer that I thought should work but doesn't

- define S so that:
 - its p points to T of lifetime 'a, and
 - its q points to T of lifetime 'b

```
struct S<'a, 'b> {  
  p: &'a T,  
  q: &'b T  
}
```

- define make_s so it returns S<'a, 'b> where 'a is the lifetime of its first parameter, like:

```
fn make_s(a: &'a T, b: &'b T) -> S<'a, 'b> {  
  S { p: a, q: b }  
}
```

The compiler complains

```
[E0597] Error: `b` does not live long enough
[command_36:1:1]
16 |     s = make_s(&a, &b);
    |               ^^^
    |               +--- borrowed value does not live long enough
17 | }
    | _
    | +--- `b` dropped here while still borrowed
18 | s.p.x
    | -----
    | +----- borrow later used here
```


- I don't know what is the exact spec of Rust that rejects this program, but it is apparently that Rust disallows *dereference of any struct any lifetime parameter of which is invalid at the point of dereference*
- in this example, `s : S<'a, 'b>` and one of its lifetime parameters (`'b`) is invalid at line 18

Summary

Why memory management is difficult

- every language wants to prevent *dereferencing a pointer to an already-reclaimed memory block (dangling pointer)*
- the problem would have been trivial if *you could reclaim v 's referent as soon as v goes out of scope*
- this is not the case, as *v 's referent may still be reachable from other variables when v goes out of scope*


```
let p : &T;  
{  
  let v = T{x: ...};  
  ...  
  p = &v;  
} // v never used below, but its referent is ... p.x ...
```



Why memory management is difficult

- every language wants to prevent *dereferencing a pointer to an already-reclaimed memory block (dangling pointer)*
- the problem would have been trivial if *you could reclaim v 's referent as soon as v goes out of scope*
- this is not the case, as *v 's referent may still be reachable from other variables when v goes out of scope*


```
let p : &T;  
{  
  let v = T{x: ...};  
  ...  
  p = &v;  
} // v never used below, but its referent is ... p.x ...
```



Why memory management is difficult

- every language wants to prevent *dereferencing a pointer to an already-reclaimed memory block (dangling pointer)*
- the problem would have been trivial if *you could reclaim v 's referent as soon as v goes out of scope*
- this is not the case, as *v 's referent may still be reachable from other variables when v goes out of scope*


```
let p : &T;  
{  
  let v = T{x: ...};  
  ...  
  p = &v;  
} // v never used below, but its referent is ... p.x ...
```




C vs. GC vs. Rust

- C/C++ : it's up to you
- GC : if it is reachable from other variables, I retain it for you
- Rust : when v goes out of scope,
 1. I reclaim T_v , all data *reachable from v through owning pointers*
 2. T_v may be reachable from other variables via borrowing references, but I guarantee such references are never dereferenced


C/C++



GC




Rust




How Rust achieved it?

- say two data structures T_v rooted at variable v and T_p rooted at variable p
- assume v goes out of scope earlier than p
- we wish to guarantee when v goes out of scope, it is safe to reclaim the entire T_v
- generally it is of course not the case, as there may be pointers somewhere in $T_p \rightarrow$ somewhere in T_v




How Rust achieved it?

- say two data structures T_v rooted at variable v and T_p rooted at variable p
- assume v goes out of scope earlier than p
- we wish to guarantee when v goes out of scope, it is safe to reclaim the entire T_v
- generally it is of course not the case, as there may be pointers somewhere in $T_p \rightarrow$ somewhere in T_v




How Rust achieved it?

- recall the “single-owner rule,” which guarantees there is only one owning pointer to any node
- \Rightarrow there can be no *owning* pointers from outside T_v to inside T_v
- \Rightarrow any such pointer must be a borrowing pointer
- crucially, a borrowing pointer must have a lifetime parameter (lifetime of the referent); say 'a




How Rust achieved it?

- recall the “single-owner rule,” which guarantees there is only one owning pointer to any node
- \Rightarrow there can be no *owning* pointers from outside T_v to inside T_v
- \Rightarrow any such pointer must be a borrowing pointer
- crucially, a borrowing pointer must have a lifetime parameter (lifetime of the referent); say 'a




How Rust achieved it?

- recall the “single-owner rule,” which guarantees there is only one owning pointer to any node
- \Rightarrow there can be no *owning* pointers from outside T_v to inside T_v
- \Rightarrow any such pointer must be a borrowing pointer
- crucially, a borrowing pointer must have a lifetime parameter (lifetime of the referent); say 'a



How Rust achieved it?

- any structure containing borrowing pointers must carry these parameters too, as part of its type (e.g., $S<'a>$)
- assignment to such borrowing pointers determines $'a$ to end when the righthand side goes out of scope (α in the figure)
- by $'a = \alpha$, the containing data structure (T_p , of type $S<'a>$) cannot be dereferenced



How Rust achieved it?

- any structure containing borrowing pointers must carry these parameters too, as part of its type (e.g., $S<'a>$)
- assignment to such borrowing pointers determines $'a$ to end when the righthand side goes out of scope (α in the figure)
- by $'a = \alpha$, the containing data structure (T_p , of type $S<'a>$) cannot be dereferenced

